

Lecture 3

CMPEN 331

A decorative blue graphic consisting of a curved line that starts near the top left and sweeps downwards and to the right, forming a large, solid blue area in the bottom right corner of the slide.

Instructions: Language of the Computer

CHAPTER 2

Instruction Set

- The repertoire of instructions of a computer
- Different computers have different instruction sets
 - But with many aspects in common
- Early computers had very simple instruction sets
 - Simplified implementation
- Many modern computers also have simple instruction sets

Introduction

- A *bit* is the most basic unit of information in a computer.
 - It is a state of “on” or “off” in a digital circuit.
 - Sometimes these states are “high” or “low” voltage instead of “on” or “off..”
- A *byte* is a group of eight bits.
 - A byte is the smallest possible *addressable* unit of computer storage.
 - The term, “addressable,” means that a particular byte can be retrieved according to its location in memory.

The MIPS Instruction Set

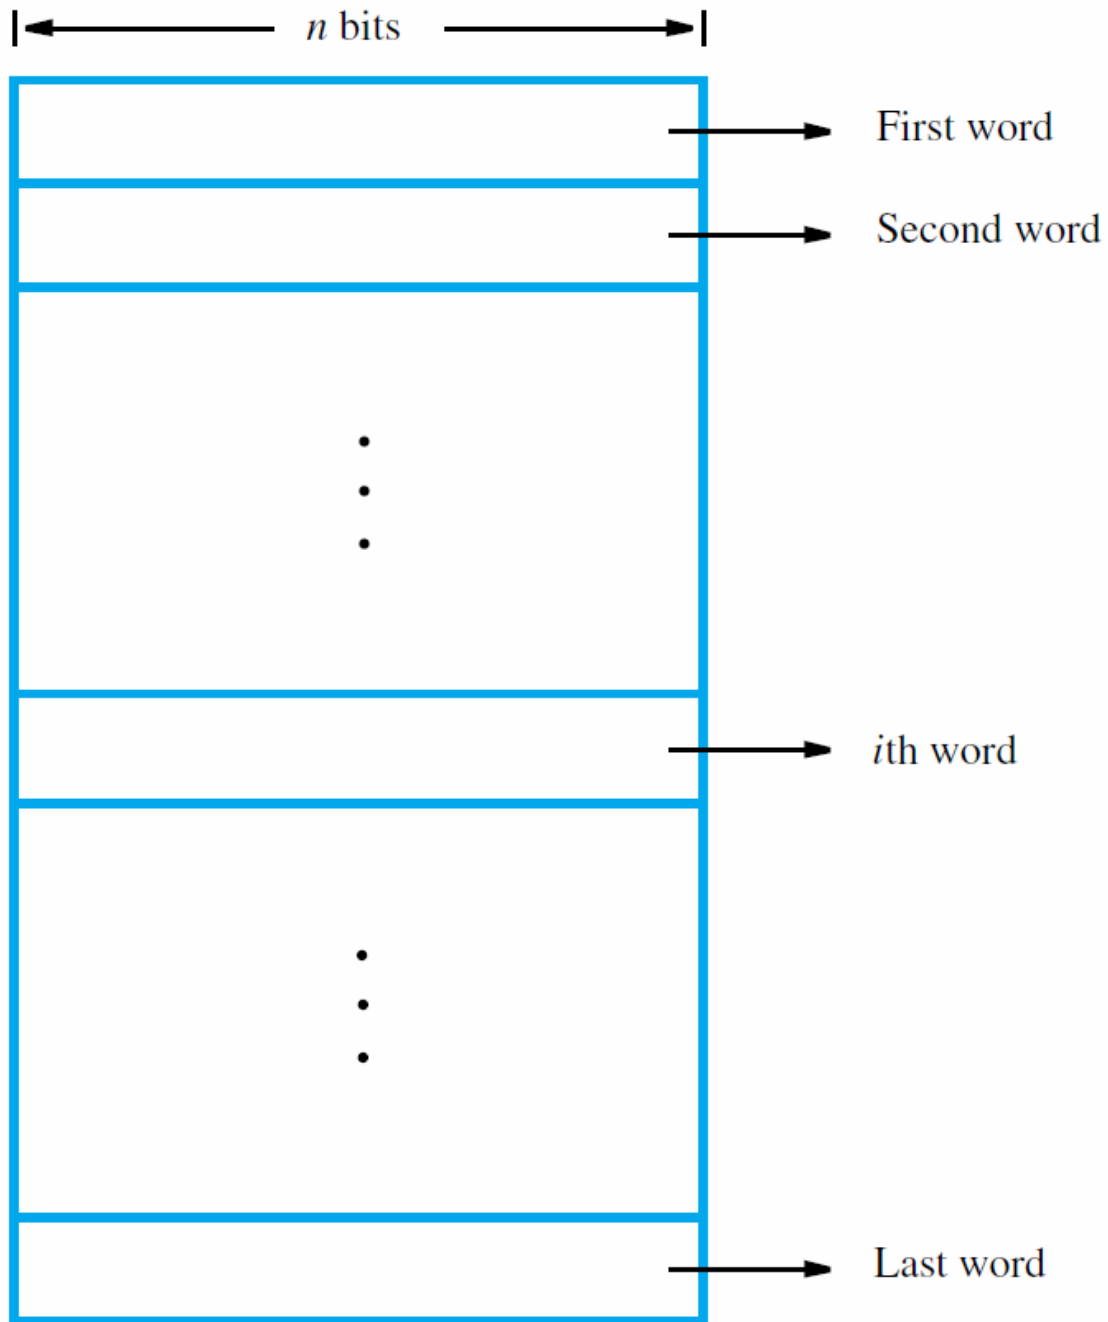
- Used as the example throughout the book
- Stanford MIPS commercialized by MIPS Technologies (www.mips.com)
- Large share of embedded core market
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- Typical of many modern ISAs

Reminders

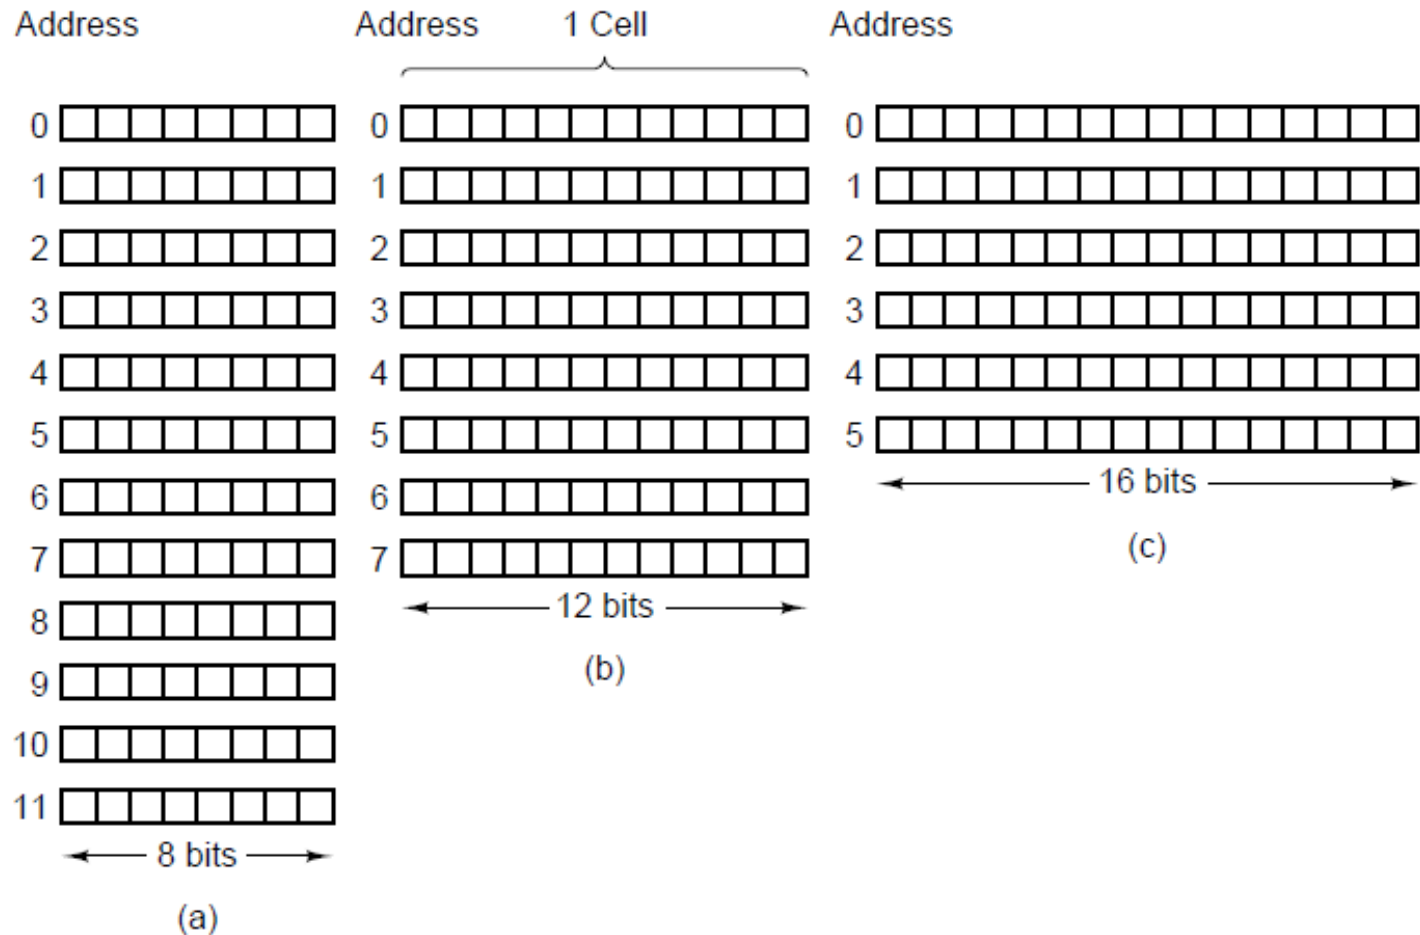
- Install *MARS* on your laptop
- Keep track of course updates on CANVAS
- Make sure your CSE account is operational
- Question/comments about the lab account/hardware/system, contact helpdesk@cse.psu.edu
- Questions about the programming assignments go to the course TAs

Memory Organization

- Memory consists of many millions of **cells**
- Each cell holds a bit of information, 0 or 1
- Information is usually handled in larger units
- A **word** is a group of n bits
- **Word length** can be 16 to 64 bits but we will stick with 32 bits as a one word through out the semester
- **Memory** is a collection of consecutive words of the size specified by the word length




Primary Memory (2)



Three ways of organizing a 96-bit memory

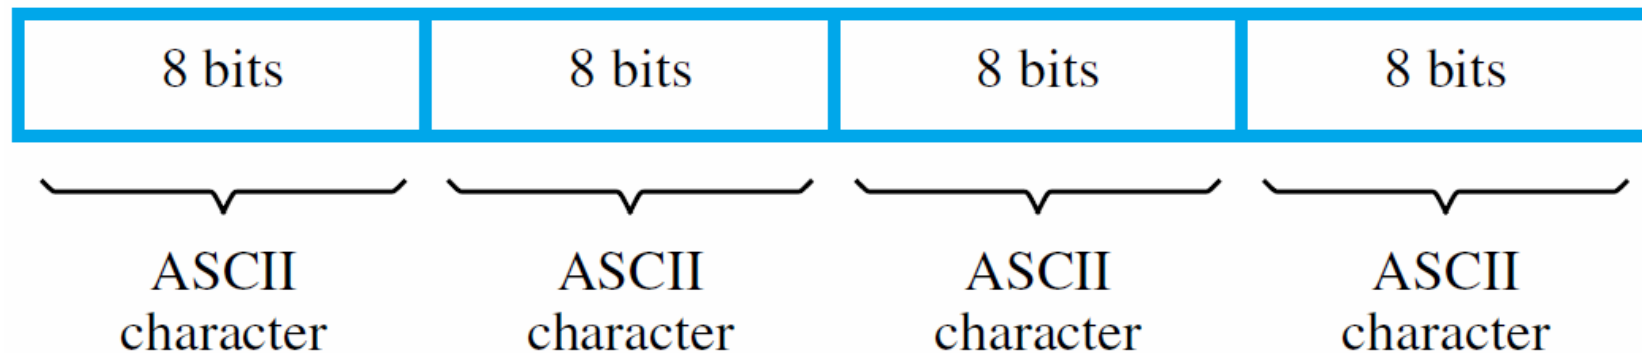
Word and Byte Encoding

- A common word length is 32 bits
- Such a word can store a 32-bit signed integer or four 8-bit bytes (e.g., ASCII characters)
- For 32-bit integer encoding, bit b_{31} is sign bit
- Words in memory may store data or machine instructions for a program
- Each machine instruction may require one or more consecutive words for encoding



↑ Sign bit: $b_{31} = 0$ for positive numbers
 $b_{31} = 1$ for negative numbers

(a) A signed integer



(b) Four characters

Addresses for Memory Locations

- To store or retrieve items of information, each memory location has a distinct **address**
- Numbers 0 to $2^k - 1$ are used as addresses for successive locations in the memory
- The 2^k locations constitute the **address space**
- Memory size set by k (number of address bits)
- Examples: $k = 20 \rightarrow 2^{20}$ or 1M locations,
 $k = 32 \rightarrow 2^{32}$ or 4G locations

Byte Addressability

- Byte size is always 8 bits
- Impractical to assign an address to each bit
- Instead, provide a **byte-addressable** memory that assigns an address to each byte
- Byte locations have addresses 0, 1, 2, ...
- Assuming that the word length is **32** bits, word locations have addresses 0, 4, 8, ...

Big- and Little-Endian Addressing

- Byte ordering, or *endianness*, is another major architectural consideration.
- Two ways to assign byte address across words
- **Big-endian** addressing assigns lower addresses to more significant (leftmost) bytes of word
- **Little-endian** addressing uses opposite order
- Addresses for 32-bit words are still 0, 4, 8, ...
- Bits in each byte labeled $b_7 \dots b_0$, left to right

Big- and Little-Endian Addressing

- If we have a two-byte integer, the integer may be stored so that the least significant byte is followed by the most significant byte or vice versa.
 - In *little endian* machines, the least significant byte is followed by the most significant byte.
 - *Big endian* machines store the most significant byte first (at the lower address).

Word
address

Byte address

0	0	1	2	3
4	4	5	6	7
$2^k - 4$	$2^k - 4$	$2^k - 3$	$2^k - 2$	$2^k - 1$

(a) Big-endian assignment

Byte address

0	3	2	1	0
4	7	6	5	4
$2^k - 4$	$2^k - 1$	$2^k - 2$	$2^k - 3$	$2^k - 4$

(b) Little-endian assignment

Big- and Little-Endian Addressing

- As an example, suppose we have the hexadecimal number 0x12345678.
- The big endian and small endian arrangements of the bytes are shown below.

Address →	00	01	10	11
Big Endian	12	34	56	78
Little Endian	78	56	34	12

Little Endian Order

- All data types larger than a byte store their individual bytes in reverse order. The least significant byte occurs at the first (lowest) memory address.

- Example:

12345678h

0000:	78
0001:	56
0002:	34
0003:	12

Word Alignment

- # of bytes per word is normally a power of 2
- Word locations have **aligned** addresses if they begin at byte addresses that are multiples of the number of bytes in a word
- Examples of aligned addresses:
 - 2 bytes per word $\rightarrow 0, 2, 4, \dots$
 - 8 bytes per word $\rightarrow 0, 8, 16, \dots$

Memory Models

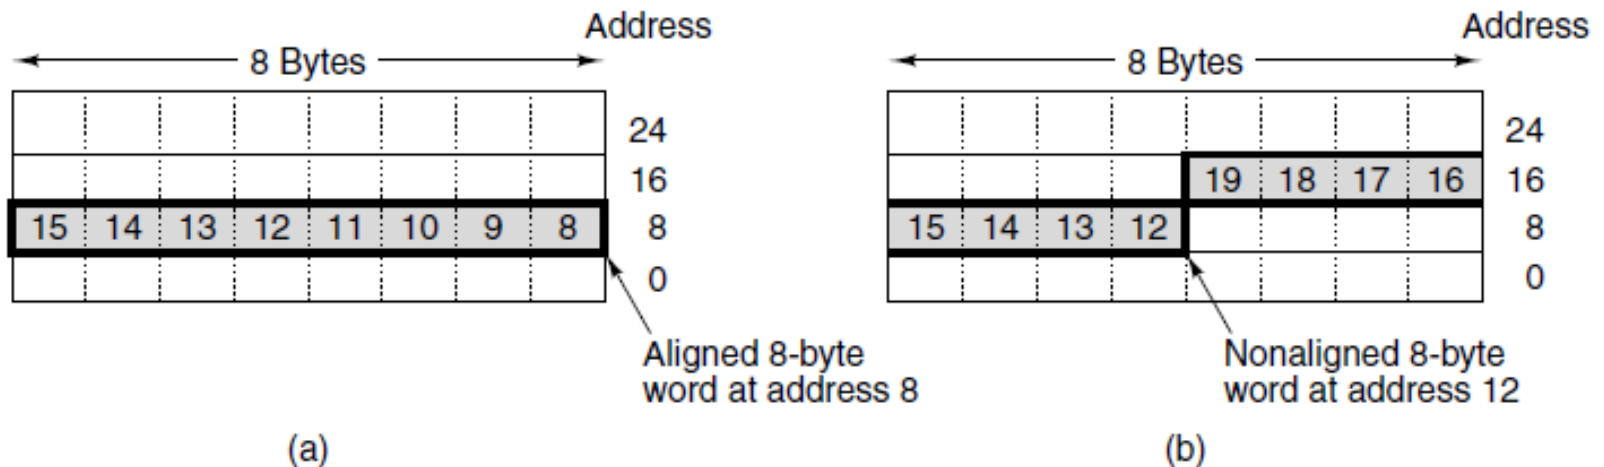


Figure 5-2. An 8-byte word in a memory.
(a) Aligned. (b) Not aligned. Some machines require that words in memory be aligned.

Memory Operations

- Memory contains data & program instructions
- Control circuits initiate transfer of data and instructions between memory and processor
- **Read operation:** memory retrieves contents at address location given by processor
- **Write operation:** memory overwrites contents at given location with given data

Instructions and Sequencing

- Instructions for a computer must support:
 - data transfers to and from the memory
 - arithmetic and logic operations on data
 - program sequencing and control
 - input/output transfers

Register Transfer Notation

- Register transfer notation is used to describe hardware-level data transfers and operations
- Arbitrary names for locations in memory
- Use [...] to denote contents of a location
- Use \leftarrow to denote transfer to a destination
- Example: $R2 \leftarrow [LOC]$
(transfer from LOC in memory to register R2)

Register Transfer Notation

- RTN can be extended to also show arithmetic operations involving locations
- Example: $R4 \leftarrow [R2] + [R3]$
(add the contents of registers R2 and R3,
place the sum in register R4)
- Right-hand expression always denotes a value, left-hand side always names a location

Assembly-Language Notation

- RTN shows data transfers and arithmetic
- Another notation is needed to represent machine instructions & programs using them
- **Assembly language** is used for this purpose
- For the two preceding examples using RTN, the assembly-language instructions are:

lw R2, LOC

add R4, R2, R3

RISC and CISC Instruction Sets

- Nature of instructions distinguishes computer
- Two fundamentally different approaches
- **Reduced Instruction Set Computers (RISC)** have one-word instructions and require arithmetic operands to be in registers
- **Complex Instruction Set Computers (CISC)** have multi-word instructions and allow operands directly from memory

RISC Instruction Sets

- RISC is simpler
- RISC instructions each occupy a single word
- A **load/store architecture** is used, meaning:
 - only Load and Store instructions are used to access memory operands
 - operands for arithmetic/logic instructions must be in registers, or one of them may be given explicitly in instruction word

Real World Architectures

- The MIPS family of CPUs has been one of the most successful in its class.
- In 1986 the first MIPS CPU was announced.
- It had a 32-bit word size and could address 4GB of memory.
- Over the years, MIPS processors have been used in general purpose computers as well as in games.
- The MIPS architecture now offers 32- and 64-bit versions.

Real World Architectures

- MIPS was one of the first RISC microprocessors.
- MIPS was designed with performance in mind: It is a *load/store* architecture, meaning that only the **load and store** instructions can access memory.
- The registers in the MIPS architecture keeps bus traffic to a minimum.

Arithmetic Operations

- Add and subtract, three operands
 - Two sources and one destination

add a, b, c # a gets b + c
- All arithmetic operations have this form
- *Design Principle 1*: Simplicity favours regularity
 - Regularity makes implementation simpler
 - Simplicity enables higher performance at lower cost

Arithmetic Example

- C code:

$f = (g + h) - (i + j);$

- Compiled MIPS code:

```
add t0, g, h    # temp t0 = g + h
add t1, i, j    # temp t1 = i + j
sub f, t0, t1   # f = t0 - t1
```

Register Operands

- Arithmetic instructions use register operands
- MIPS has a 32×32 -bit register file
 - Use for frequently accessed data
 - Numbered 0 to 31
 - 32-bit data called a “word”
- Assembler names
 - \$t0, \$t1, ..., \$t9 for temporary values
 - \$s0, \$s1, ..., \$s7 for saved variables
- *Design Principle 2: Smaller is faster*
 - main memory: millions of locations

Register Operand Example

- C code:

`f = (g + h) - (i + j);`

- `f, ..., j` in `$s0, ..., $s4`

- Compiled MIPS code:

`add $t0, $s1, $s2`

`add $t1, $s3, $s4`

`sub $s0, $t0, $t1`

Memory Operands

- Main memory used for composite data
 - Arrays, structures, dynamic data
- To apply arithmetic operations
 - Load values from memory into registers
 - Store result from register to memory
- Memory is byte addressed
 - Each address identifies an 8-bit byte
- Words are aligned in memory
 - Address must be a multiple of 4
- MIPS is **Big Endian**
 - Most-significant byte at least address of a word

Memory Operand Example 1

- C code:

```
g = h + A[8];
```

- g in \$s1, h in \$s2, base address of A in \$s3

- Compiled MIPS code:

- Index 8 requires offset of 32
 - 4 bytes per word

```
lw    $t0, 32($s3)    # load word  
add   $s1, $s2, $t0
```

offset



base register



Lecture 4

CMPEN 331

A decorative blue graphic consisting of a curved line that starts near the top left and sweeps down towards the bottom right, forming a large, stylized shape that resembles a quarter-circle or a wedge.

Memory Operand Example 2

- C code:

`A[12] = h + A[8];`

- `h` in `$s2`, base address of `A` in `$s3`

- Compiled MIPS code:

- Index 8 requires offset of 32

```
lw    $t0, 32($s3)    # load word
add   $t0, $s2, $t0
sw    $t0, 48($s3)    # store word
```

Registers vs. Memory

- Registers are faster to access than memory
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler must use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Immediate Operands

- Constant data specified in an instruction, add immediate
`addi $s3, $s3, 4`
- No subtract immediate instruction
 - Just use a negative constant
`addi $s2, $s1, -1`
- *Design Principle 3: Make the common case fast*
 - Small constants are common
 - Immediate operand avoids a load instruction

The Constant Zero

- MIPS register 0 (\$zero) is the constant 0
 - Cannot be overwritten
- Useful for common operations
 - E.g., move between registers
add \$t2, \$s1, \$zero

Unsigned Binary Integers

- Given an n-bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$
- Example
 - 0000 0000 0000 0000 0000 0000 0000 1011₂
= $0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
= $0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$
- Using 32 bits
 - 0 to +4,294,967,295

2s-Complement Signed Integers

- Given an n-bit number

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$

- Example

- $1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2$
 $= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0$
 $= -2,147,483,648 + 2,147,483,644 = -4_{10}$

- Using 32 bits

- $-2,147,483,648$ to $+2,147,483,647$

2s-Complement Signed Integers

- Bit 31 is sign bit
 - 1 for negative numbers
 - 0 for non-negative numbers
- Non-negative numbers have the same unsigned and 2s-complement representation
- Some specific numbers
 - 0: 0000 0000 ... 0000
 - -1: 1111 1111 ... 1111
 - Most-negative: 1000 0000 ... 0000
 - Most-positive: 0111 1111 ... 1111

Signed Negation

- Complement and add 1
 - Complement means $1 \rightarrow 0, 0 \rightarrow 1$

$$x + \bar{x} = 1111 \dots 111_2 = -1$$

$$\bar{x} + 1 = -x$$

- Example: negate +2
 - $+2 = 0000\ 0000 \dots 0010_2$
 - $-2 = 1111\ 1111 \dots 1101_2 + 1$
 $= 1111\ 1111 \dots 1110_2$

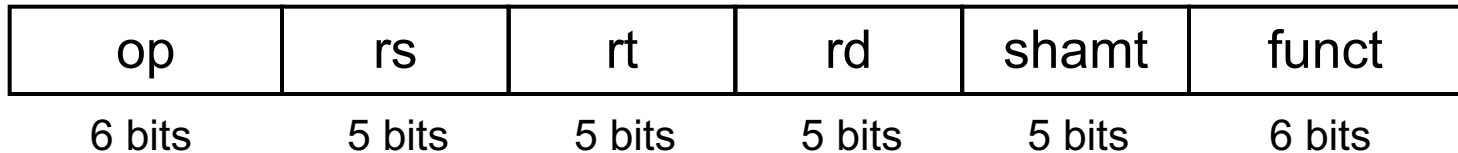
Sign Extension

- Representing a number using more bits
 - Preserve the numeric value
- In MIPS instruction set
 - `addi`: extend immediate value
 - `lb`, `lh`: extend loaded byte/halfword
- Replicate the sign bit to the left
 - unsigned values: extend with 0s
- Examples: 8-bit to 16-bit
 - +2: 0000 0010 => 0000 0000 0000 0010
 - -2: 1111 1110 => 1111 1111 1111 1110

Representing Instructions

- Instructions are encoded in binary
 - Called machine code
- MIPS instructions
 - Encoded as 32-bit instruction words
 - Small number of formats encoding operation code (opcode), register numbers, ...
 - Regularity!
- Register numbers
 - \$t0 – \$t7 are reg's 8 – 15
 - \$t8 – \$t9 are reg's 24 – 25
 - \$s0 – \$s7 are reg's 16 – 23

MIPS R-format Instructions



- Instruction fields
 - op: operation code (opcode)
 - rs: first source register number
 - rt: second source register number
 - rd: destination register number
 - shamt: shift amount (00000 for now)
 - funct: function code (extends opcode)

R-format Example

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

add \$t0, \$s1, \$s2

special	\$s1	\$s2	\$t0	0	add
0	17	18	8	0	32
000000	10001	10010	01000	00000	100000

$00000010001100100100000000100000_2 = 02324020_{16}$