



PennState

CMPSC 311 - Introduction to Systems Programming

Introduction to C


Professors:

Sencun Zhu and Suman Saha

(Slides are mostly by *Professor Patrick McDaniel*
and *Professor Abutalib Aghayev*)



But first, emacs or IDE?



Emacs users



Visual Studio user

Emacs in the movies..



PennState



Emacs in the movies..




What makes emacs powerful?



- It's programmable with a full-fledged functional programming language, LISP

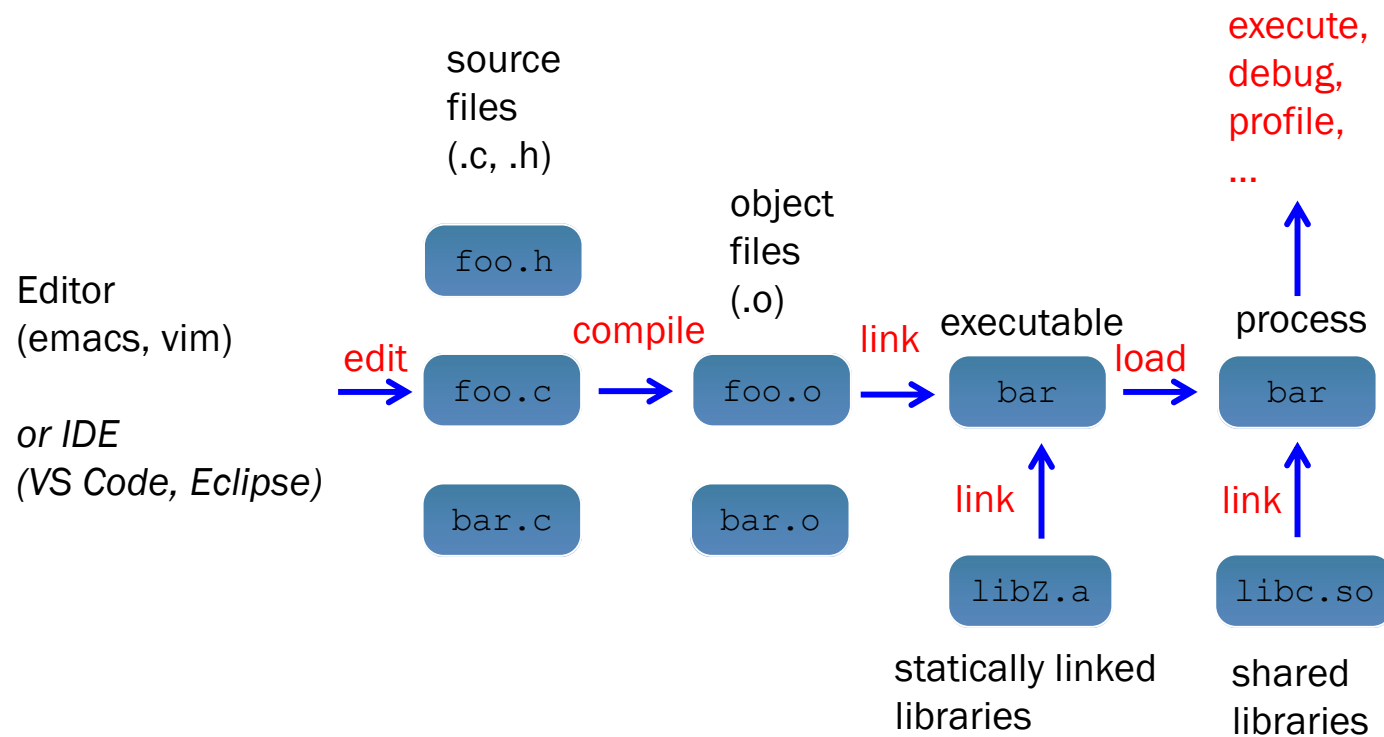
1.00



C workflow



PennState



Defining a function



PennState

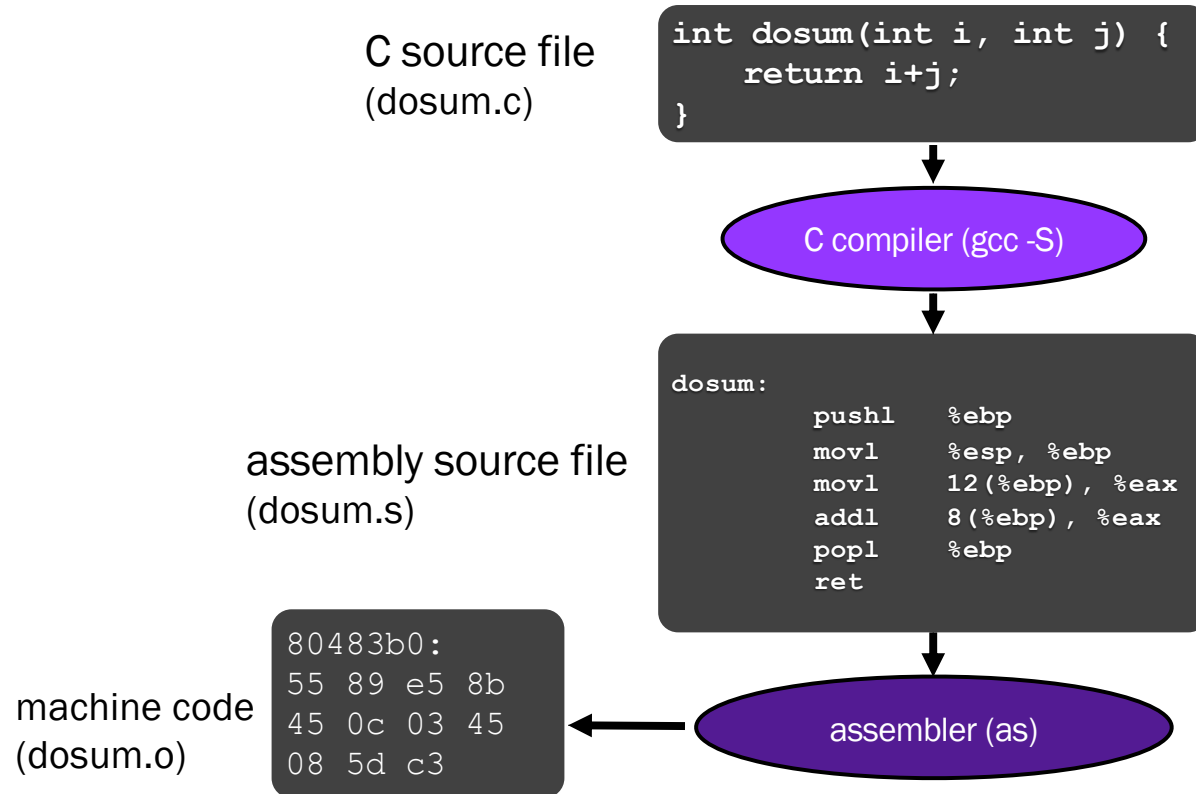
```
returnType name(type name, ..., type name)  
{  
    statements;  
}
```

```
// sum integers from 1 to max  
int sumTo(int max) {  
    int i, sum = 0;  
  
    for (i=1; i<=max; i++) {  
        sum += i;  
    }  
    return sum;  
}
```

From C to machine code




PennState





- Most C compilers generate object “.o” files directly
 - i.e., without actually saving the readable .s assembly file



Note: *Object code* is re-locatable machine code, but generally cannot be executed without some manipulation (e.g., via a linker)

Anatomy of a C program ...



PennState

```
#include <stdio.h>

int myfunc(int i) {
    printf("Got into function with %d\n", i);
    return 0;
}

int main(void) {
    myfunc(10);
    return 0;
}
```

All C programs start with the “main()” function ...

Anatomy of a C program ...



PennState

```
#include <stdio.h>

int myfunc(int i) {
    printf("Got into function with %d\n", i);
    return 0;
}

int main(void) {
    myfunc(10);
    return 0;
}
```

Compile and link 

```
%_gcc -g -Wall main.c -o main
%_./main
Got into function with 10
%_
```

Running the program 

Running a program



```
mcdaniel@ubuntu:~/tmp/helloworld$ emacs helloworld.c
mcdaniel@ubuntu:~/tmp/helloworld$ gcc helloworld.c -o helloworld
mcdaniel@ubuntu:~/tmp/helloworld$ helloworld
helloworld: command not found
mcdaniel@ubuntu:~/tmp/helloworld$ echo $PATH
/usr/lib/lightdm/lightdm:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin
:/sbin:/bin:/usr/games:/usr/local/games
mcdaniel@ubuntu:~/tmp/helloworld$ ./helloworld
Hello world!
mcdaniel@ubuntu:~/tmp/helloworld
```

- UNIX looks for a program in all of the directories listed by the **PATH** environment variable, or locally if prepended by “**./**”
 - to add to search path just add more “**:**” separated paths,

```
export PATH=$PATH:/new/path
```

Multi-file C programs



PennState

C source file
(dosum.c)

```
int dosum(int i, int j) {  
    return i+j;  
}
```

this “prototype” of
dosum() tells gcc
about the types of
dosum’s arguments
and its return value

C source file
(sumnum.c)

```
#include <stdio.h>  
int dosum(int i, int j);  
  
int main(int argc, char **argv)  
{  
    printf("%d\n", dosum(1,2));  
    return 0;  
}
```

dosum() is
implemented
in dosum.c

Multi-file C programs



C source file
(dosum.c)

```
int dosum(int i, int j) {  
    return i+j;  
}
```

C source file
(sumnum.c)

```
#include <stdio.h>  
  
int dosum(int i, int j);  
  
int main(int argc, char **argv)  
{  
    printf("%d\n", dosum(1,2));  
    return 0;  
}
```


why do we need this
#include?

where is the
implementation
of printf?

Compiling multi-file programs



- Multiple object files are *linked* to produce an executable
 - standard libraries (libc, crt1, ...) are usually also linked in
 - a library is just a pre-assembled collection of .o files



Object files revisited ...



PennState

- `sumnum.o`, `dosum.o` are **object files**
 - each contains machine code produced by the compiler
 - each might contain references to external symbols
 - variables and functions not defined in the associated `.c` file
 - e.g., `sumnum.o` contains code that relies on `printf()` and `dosum()`, but these are defined in `libc.a` and `dosum.o`, respectively
- linking **resolves** these external symbols while smooshing together object files and libraries

Lets dive into C itself



PennState

- Things that are the same as Java
- syntax for statements, control structures, function calls
- types: `int, double, char, long, float`
- type-casting syntax: `float x = (float) 5 / 3;`
- expressions, operators, precedence

`+ - * / % ++ -- = += -= *= /= %= < <= == != > >= && || !`

- scope (local scope is within a set of `{ }` braces)
- comments: `/* comment */` or `// comment *to EOL*`

Primitive types in C



- integer types
 - `char`, `int`
- floating point
 - `float`, `double`
- modifiers
 - `short [int]`
 - `long [int, double]`
 - `signed [char, int]`
 - `unsigned [char, int]`

type	bytes (32-bit)	bytes (64-bit)	32-bit range	printf
char	1	1	[0, 255]	%c
short int	2	2	[-32768, 32767]	%hd
unsigned short int	2	2	[0, 65535]	%hu
int	4	4	[-214748648, 2147483647]	%d
unsigned int	4	4	[0, 4294967295]	%u
long int	4	8	[-2147483648, 2147483647]	%ld
long long int	8	8	[-9223372036854775808, 9223372036854775807]	%lld
float	4	4	approx [10 ⁻³⁸ , 10 ³⁸]	%f
double	8	8	approx [10 ⁻³⁰⁸ , 10 ³⁰⁸]	%lf
long double	12	16	approx [10 ⁻⁴⁹³² , 10 ⁴⁹³²]	%Lf
pointer	4	8	[0, 4294967295]	%p

C99 extended integer types



PennState

- Solve the conundrum of “how big is a long int?”

```
#include <stdint.h>

void foo(void) {
    int8_t  w;    // exactly 8 bits, signed
    int16_t x;    // exactly 16 bits, signed
    int32_t y;    // exactly 32 bits, signed
    int64_t z;    // exactly 64 bits, signed

    uint8_t w;    // exactly 8 bits, unsigned
    ...etc.
}
```

Similar to Java...



- variables
 - must declare at the start of a function or block (not required since in C99)
 - need not be initialized before use (gcc -Wall will warn); **ALWAYS INITIALIZE YOUR VARS**

```
#include <stdio.h>

int main(void) {
    int x, y = 5;    // note x is uninitialized!
    long z = x+y;

    printf("z is '%ld'\n", z); // what's printed?
    {
        int y = 10;
        printf("y is '%d'\n", y);
    }
    int w = 20;    // ok in c99
    printf("y is '%d', w is '%d'\n", y, w);
    return 0;
}
```

Similar to Java...



PennState

- `const`
 - a qualifier that indicates the variable's value cannot change
 - compiler will issue an error if you try to violate this
 - why is this qualifier useful?

```
#include <stdio.h>

int main(void) {
    const double MAX_GPA = 4.0;

    printf("MAX_GPA: %g\n", MAX_GPA);
    MAX_GPA = 5.0; // illegal!
    return 0;
}
```

Similar to Java...



PennState

- for loops
 - can't declare variables in the loop header (changed in c99)
- if/else, while, and do while loops
 - no boolean type (changed in c99: `#include <stdbool.h>`)
 - any type can be used; 0 means **false**, everything else **true**

```
int i;

for (i=0; i < 100; i++) {
    if (i % 10 == 0) {
        printf("i: %d\n", i);
    }
}
```

Pointers




PennState

```
#include <stdio.h>

int main(void) {
    int i = 5;
    int *ip = &i;

    printf("%d\n", i);
    printf("%p\n", ip);
    *ip = 42;
    printf("%d\n", i);
    printf("%d\n", *ip);
}
```



Key concepts:

- Taking address of a variable: `&`
- Dereferencing a pointer: `*`
- Aliasing: `*ip` is an alias for `i`

Similar to Java



PennState

- C always **passes arguments by value**
 - value is “copied” into function
 - any local modification change is not reflecting in original value passed
- pointers let you **pass by reference**
 - pass “memory location” of variable
 - more on these soon
 - least intuitive part of C
 - very dangerous part of C


```
void add_pbv(int c) {  
    c += 10;  
    printf("pbv c: %d\n", c);  
}  
  
void add_pbr(int *c) {  
    *c += 10;  
    printf("pbr *c: %d\n", *c);  
}  
  
int main(void) {  
    int x = 1;  
    printf("x: %d\n", x);  
    add_pbv(x);  
    printf("x: %d\n", x);  
    add_pbr(&x);  
    printf("x: %d\n", x);  
    return 0;  
}
```


Pass-by-value



PennState

- C passes arguments by **value**
 - callee receives a **copy** of the argument




```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

- if the callee (function that is called) modifies an argument, caller's copy isn't modified

Pass-by-value



PennState

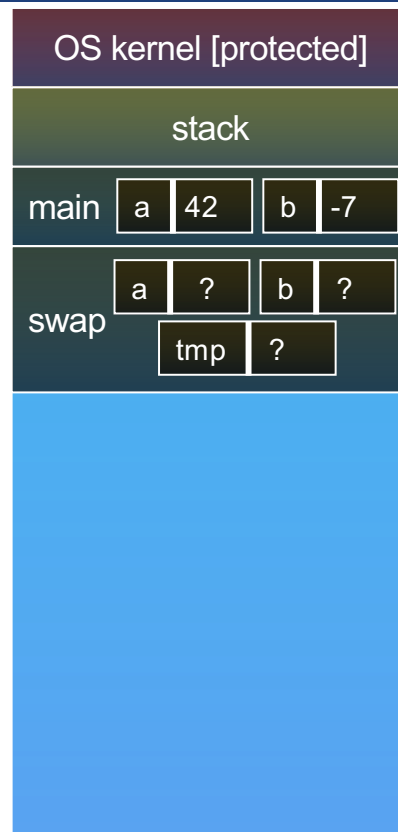


```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-value



PennState



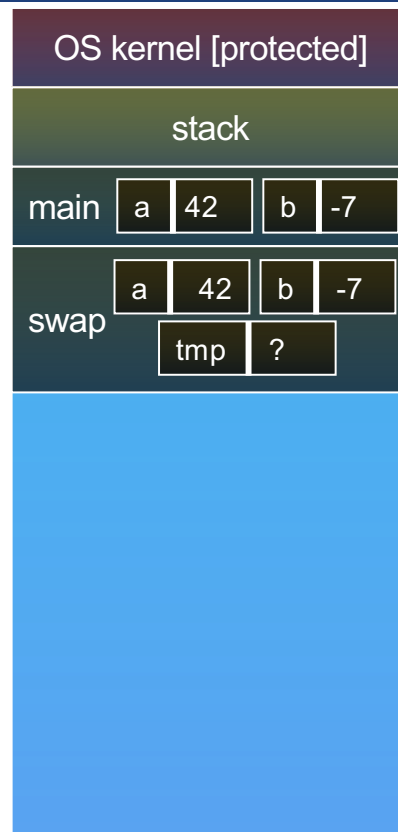
```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```



Pass-by-value



PennState

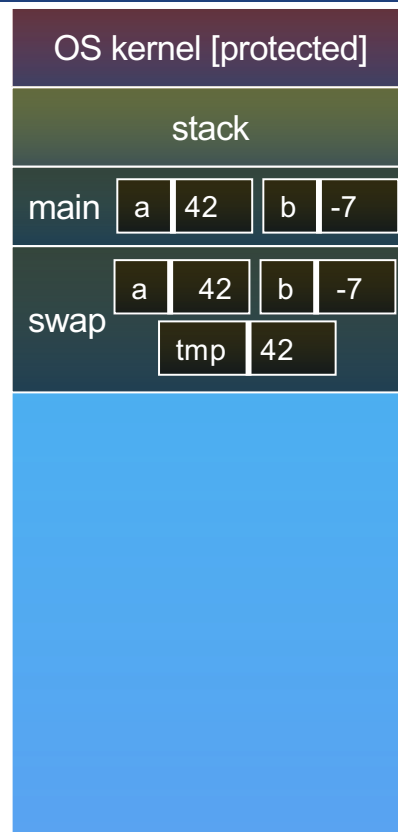


```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-value



PennState

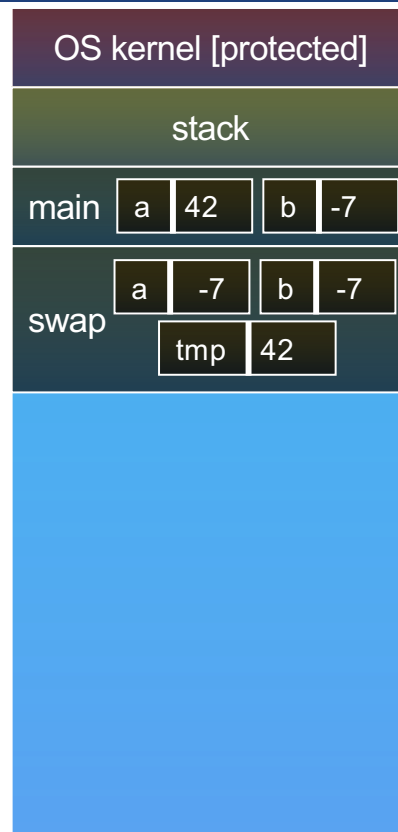


```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-value



PennState

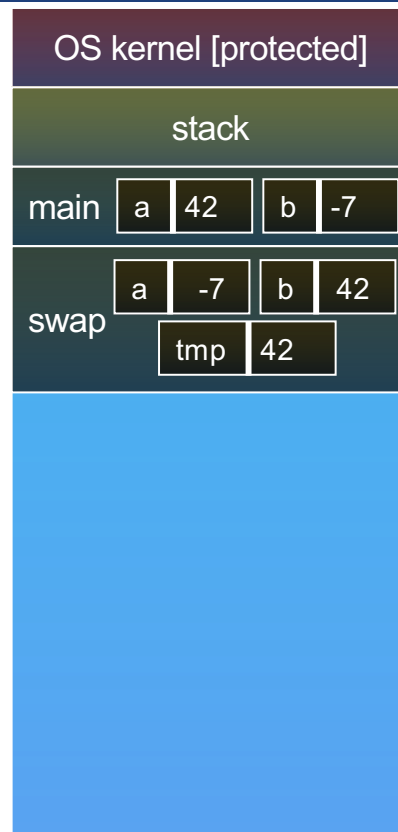


```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-value



PennState




```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-value



PennState




```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```



Pass-by-value



PennState



```
void swap(int a, int b) {  
    int tmp = a;  
    a = b;  
    b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(a, b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```




Pass-by-reference



PennState

- You can use pointers to pass by **reference**
 - callee still receives a **copy** of the argument
 - but, the argument is a **pointer***
 - the pointer's value points-to the variable in the scope of the caller
 - this gives the callee a way to modify a variable that's in the scope of the caller




```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

* The key to C (and languages like it) is getting good at using pointers.

Pass-by-reference



PennState

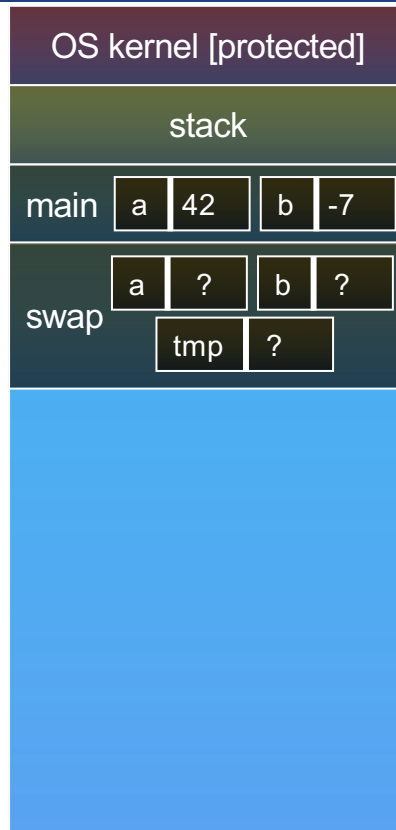


```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference

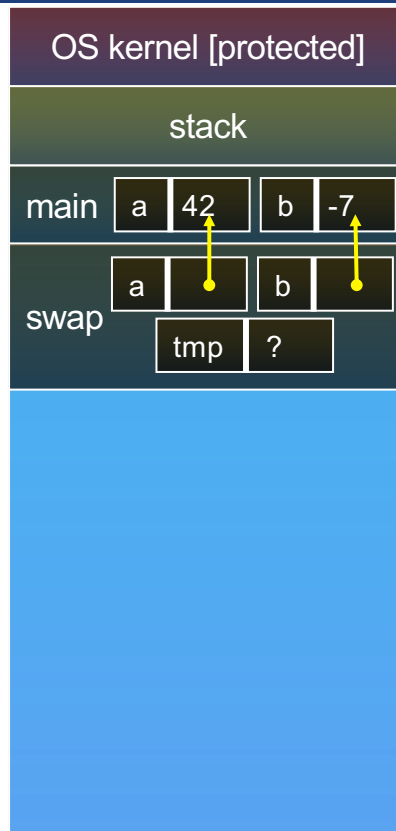


PennState



```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference

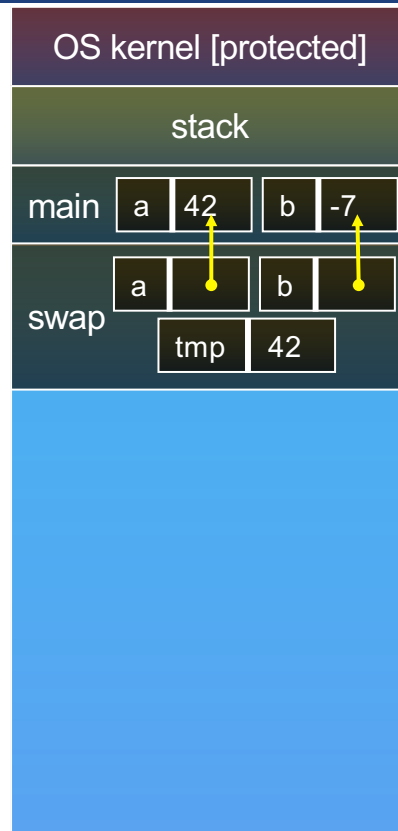


```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference



PennState

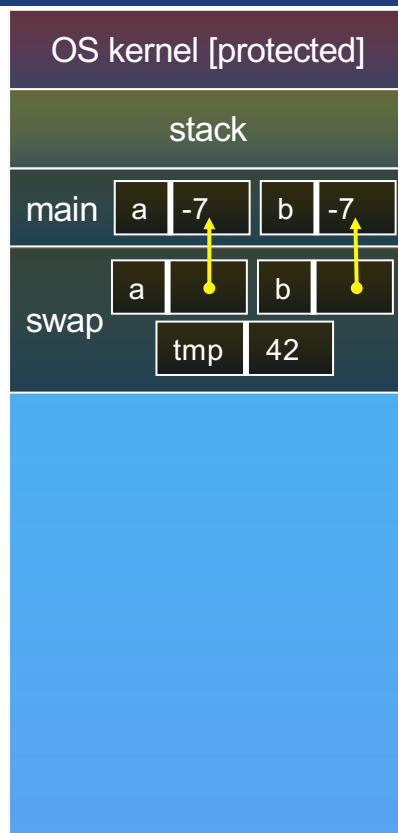


```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference



PennState

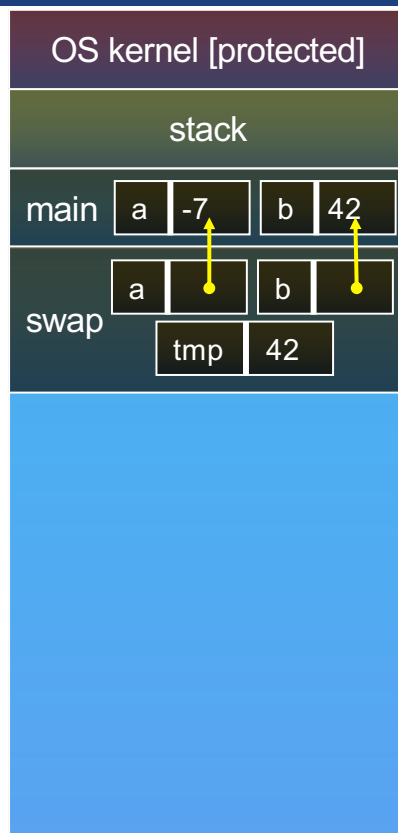


```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference



PennState




```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```


Pass-by-reference



PennState




```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```

Pass-by-reference



PennState



```
void swap(int *a, int *b) {  
    int tmp = *a;  
    *a = *b;  
    *b = tmp;  
}  
  
int main(void) {  
    int a = 42, b = -7;  
  
    swap(&a, &b);  
    printf("a: %d, b: %d\n", a, b);  
    return 0;  
}
```



Very different than Java



- arrays
 - just a bare, contiguous block of memory of the correct size
 - array of 6 integers requires 6×4 bytes = 24 bytes of memory
- arrays have no methods, do not know their own length (no **bounds checking**)
 - C doesn't stop you from overstepping the end of an array!
 - many, many security bugs come from this (buffer overflow)

```
int x[6];
```

A[0]	A[1]	A[2]	A[3]	A[4]	A[5]
34	11	-129	49	708	-11

Very different than Java



- arrays
 - just a bare, contiguous block of memory of the correct size
 - array of 6 integers requires 6×4 bytes = 24 bytes of memory
- arrays have no methods, do not know their own length (no bounds checking)
 - C doesn't **`x[7] = 45; // Legal C, but can cause memory fault!!!!`**
 - many, many security bugs come from this (buffer overflow)

```
int x[6];
```

A[0]	A[1]	A[2]	A[3]	A[4]	A[5]
34	11	-129	49	708	-11

Very different than Java



- strings
 - array of **char**
 - terminated by the **NULL** character '**\0**'
 - are not objects, have no methods; **string.h** has helpful utilities (see strings lecture coming soon!)



```
char *x = "hello\n";
```

Very different than Java



- errors and exceptions
 - C has no exceptions (no try / catch)
 - errors are returned as integer error codes from functions
 - sometimes makes error handling ugly and inelegant
 - some support from OS using signals (end of semester)
- crashes
 - if you do something bad, you'll end up spraying bytes around memory
 - hopefully causing a “segmentation fault” and crash
- objects
 - there aren't any; struct is closest feature (set of fields)



Very different than Java



- memory management
 - there is no garbage collector
 - anything you allocate you have to free ([memory leaks](#))
 - local variables are allocated off of the stack
 - freed when you return from the function
 - global and static variables are allocated in a data segment
 - are freed when your program exits
 - you can allocate memory in the heap segment using `malloc()`
 - you must free malloc'ed memory with `free()`
 - failing to free is a leak, [double-freeing](#) is an error (hopefully crash)

Very different than Java



PennState

- console I/O
 - C library (`libc`) has portable routines for reading/writing, e.g., `scanf()`, `printf()`
- file I/O
 - C library has portable routines for reading/writing
 - `fopen()`, `fread()`, `fwrite()`, `fclose()`, etc.
 - does `buffering` by default, is `blocking` by default
 - OS provides system calls
 - we'll be using these: more control over buffering, blocking
 - Low level binary reads and writes, e.g., `read()`, `write()`, `open()`, `close()`

Very different than Java



- network I/O
 - C standard library has no notion of network I/O
 - OS provides (somewhat portable) routines
 - lots of complexity lies here
 - errors: network can fail
 - performance: network can be slow
 - concurrency: servers speak to thousands of clients simultaneously

Note: most of these topics will be covered in detail over the semester.

Very different than Java



- Libraries you can count on
 - C has very few compared to most other languages
 - no built-in trees, hash tables, linked lists, sort , etc.
 - you have to write many things on your own
 - particularly data structures
 - error prone, tedious, hard to build efficiently and portably
 - less productive language than Java, C++, python, or others



Problem: ordering



PennState

- Don't call a function that hasn't been declared yet:

```
#include <stdio.h>

int main(void) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}

// sum integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;

    for (i=1; i<=max; i++) {
        sum += i;
    }
    return sum;
}
```

Problem: ordering



PennState

- Solution 1: reverse order of definition

```
#include <stdio.h>

// sum integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;

    for (i=1; i<=max; i++) {
        sum += i;
    }
    return sum;
}

int main(void) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}
```

Problem: ordering



PennState

- Solution 2: provide function declaration
 - teaches the compiler the argument and return types of the function that will appear later
- The body-less function declaration is called a *functional prototype*.

```
#include <stdio.h>

// this function prototype is a
// declaration of sumTo
int sumTo(int);

int main(void) {
    printf("sumTo(5) is: %d\n", sumTo(5));
    return 0;
}

// sum integers from 1 to max
int sumTo(int max) {
    int i, sum = 0;

    for (i=1; i<=max; i++) {
        sum += i;
    }
    return sum;
}
```



- There are three predefined streams provided to all UNIX programs
 - Standard input (`stdin`)
 - Standard output (`stdout`)
 - Standard error (`stderr`)
- `printf("this is printed to standard output\n");`
- `fprintf(stdout, "this is printed to standard output as well\n");`
- `fprintf(stderr, "this is printed to standard error\n");`