



PennState

CMPSC 311 - Introduction to Systems Programming

Bit/Byte Operations

Professors

Sencun Zhu and Suman Saha

(Slides are mostly by Professor Patrick McDaniel
and Professor Abutalib Aghayev)



Hobbit

Hobbyte

Number Systems



- All base-X systems have the following characteristic:

Assume a base b and digits $P = \{p_k, p_{k-1}, p_{k-2}, \dots, p_1, p_0\}$

$$value = \sum_{i=0}^k b^i * p_i$$

where $\forall p_i \in P, p_i = [0, b - 1]$

Base-10 (decimal)

- Digits: 0,1,2,3,4,5,6,7,8,9
- Place values: $10^0, 10^1, 10^2, \dots$
- Example: $123 = 3 \times 1 + 2 \times 10 + 1 \times 100$
 $= 123$

Base-2(binary)

- Digits: 0,1
- Place values: $2^0, 2^1, 2^2, \dots$
- Example: $1011 = 1 \times 1 + 1 \times 2 + 0 \times 4 + 1 \times 8$
 $= 11$

Base-16 (hexadecimal)

- Digits: 0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F
- Place values: $16^0, 16^1, 16^2, \dots$
- Example: $0xAFC = C \times 1 + F \times 16 + A \times 256$
 $= 2812$

Converting Decimal to Binary

- while $n \neq 0$:
 - next binary digit (from right to left) = $n \% 2$
 - $n = n / 2$

| n | remainder | digit # |
|-----|-----------|---------|
| 235 | 1 | 0 |
| 117 | 1 | 1 |
| 58 | 0 | 2 |
| 29 | 1 | 3 |
| 14 | 0 | 4 |
| 7 | 1 | 5 |
| 3 | 1 | 6 |
| 1 | 1 | 7 |

235 (base-10) = 11101011 (base-2)

Converting Decimal and Binary to Hex



- Converting decimal to hexadecimal
 - while $n \neq 0$:
 - next hex digit (from right to left) = $n \% 16$
 - $n = n / 16$

| n | remainder | digit # |
|-----|-----------|---------|
| 235 | B (11) | 0 |
| 14 | E (14) | 1 |

235 (base-10) = EB (base-16)

- Converting binary to hexadecimal
 - group binary digits to into groups of 4 bits (nibbles) starting from right
 - Convert each nibble to hexadecimal digit

235 (base-10) = 11101011 (base-2) = 1110 1011 = EB

- Be familiar with hex notation: 0xDEADBEEF has 8 nibbles, 4 bytes (or octets)
- Hexadecimal to binary: just convert each nibble to binary

Conversion Summary




- Binary to decimal
 - sum of binary digits times powers of 2
- Hexadecimal to decimal
 - sum of hex digits times power of 16
- Decimal to binary
 - division method
- Decimal to hex
 - division method
 - or, convert to binary and use binary to hex method below
- Binary to hex
 - group binary digits to nibbles, convert nibbles to hex digits
- Hex to binary
 - convert each hex digit to binary



Byte Ordering Example

- How should bytes within a multi-byte word be ordered in memory?
 - Big Endian: Least significant byte has *highest* address (SPARC)
 - Little Endian: Least significant byte has *lowest* address (x86)
- Example: `int x = 305419896`
 - Variable `x` has 4-byte representation `0x12345678`
 - Address given by `&x` is `0x100`

Big Endian



Little Endian



Examining Data Representations



- Code to find endianness of the architecture
 - Casting pointer to `uint8_t *` creates byte array

```
#include <stdio.h>
#include <stdint.h>

void show_bytes(uint8_t *start, int len) {
    for (int i = 0; i < len; ++i)
        printf("%p\t0x%.2x\n", start+i, start[i]);
    printf("\n");
}

int main(void) {
    int a = 305419896;
    printf("a lives at address %p\n\n", &a);
    show_bytes((uint8_t *)&a, sizeof(int));
}
```

Result (Linux on x86):

a lives at address 0x7ffebf803174

| | |
|----------------|------|
| 0x7ffebf803174 | 0x78 |
| 0x7ffebf803175 | 0x56 |
| 0x7ffebf803176 | 0x34 |
| 0x7ffebf803177 | 0x12 |

printf directives

%p Print pointer

%x Print hexadecimal

Boolean Algebra

- Developed by George Boole in 19th Century
 - Algebraic representation of logic that based on “True” (as 1) and “False” (as 0)

And

$A \& B = 1$ when both $A=1$ and $B=1$

| $\&$ | 0 | 1 |
|------|---|---|
| 0 | 0 | 0 |
| 1 | 0 | 1 |

Or

$A | B = 1$ when either $A=1$ or $B=1$

| $ $ | 0 | 1 |
|-----|---|---|
| 0 | 0 | 1 |
| 1 | 1 | 1 |

Not

$\sim A = 1$ when $A=0$

| \sim | |
|--------|---|
| 0 | 1 |
| 1 | 0 |


Exclusive-Or (Xor)

$A \wedge B = 1$ when either $A=1$ or $B=1$, but not both

| \wedge | 0 | 1 |
|----------|---|---|
| 0 | 0 | 1 |
| 1 | 1 | 0 |

Bit-Level Operations in C

- Operations `&`, `|`, `~`, `^` Available in C
 - Apply to any “integral” data type
 - long, int, short, char, unsigned
 - View arguments as bit vectors
 - Arguments applied bit-wise
- Examples (Char data type)
 - `~0x41` \rightarrow `0xBE`
 - `~010000012` \rightarrow `101111102`
 - `~0x00` \rightarrow `0xFF`
 - `~000000002` \rightarrow `111111112`
 - `0x69 & 0x55` \rightarrow `0x41`
 - `011010012 & 010101012` \rightarrow `010000012`
 - `0x69 | 0x55` \rightarrow `0x7D`
 - `011010012 | 010101012` \rightarrow `011111012`



Contrast: Logic Operations in C

- Contrast to Logical Operators (`&&`, `||`, `!`)
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1
 - Early termination

Representing & Manipulating Sets PennState

- Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$ for set “A”
- $a_j = 1$ if $j \in A$

01010101 { 0, 2, 4, 6 }
76543210

01101001 { 0, 3, 5, 6 }
76543210

Operations On Sets:

| | | |
|------------------------|----------|----------------------|
| & Intersection | 01000001 | { 0, 6 } |
| Union | 01111101 | { 0, 2, 3, 4, 5, 6 } |
| ^ Symmetric difference | 00111100 | { 2, 3, 4, 5 } |
| ~ Complement | 10101010 | { 1, 3, 5, 7 } |

Representing Signed Numbers

| b3 | b2 | b1 | b0 | | Unsigned | One's complement | Two's complement |
|----|----|----|----|--|----------|------------------|-------------------|
| 0 | 0 | 0 | 0 | | 0 | 0 | 0 |
| 0 | 0 | 0 | 1 | | 1 | 1 | 1 |
| 0 | 0 | 1 | 0 | | 2 | 2 | 2 |
| 0 | 0 | 1 | 1 | | 3 | 3 | 3 |
| 0 | 1 | 0 | 0 | | 4 | 4 | 4 |
| 0 | 1 | 0 | 1 | | 5 | 5 | 5 |
| 0 | 1 | 1 | 0 | | 6 | 6 | 6 |
| 0 | 1 | 1 | 1 | | 7 | 7 | $7 = 2^{n-1} - 1$ |
| 1 | 0 | 0 | 0 | | 8 | -7 | $-8 = -2^{n-1}$ |
| 1 | 0 | 0 | 1 | | 9 | -6 | -7 |
| 1 | 0 | 1 | 0 | | 10 | -5 | -6 |
| 1 | 0 | 1 | 1 | | 11 | -4 | -5 |
| 1 | 1 | 0 | 0 | | 12 | -3 | -4 |
| 1 | 1 | 0 | 1 | | 13 | -2 | -3 |
| 1 | 1 | 1 | 0 | | 14 | -1 | -2 |
| 1 | 1 | 1 | 1 | | 15 | -0 | -1 |

- Computing one's complement negative representation:
 - Complement the positive number
- Computing two's complement negative representation:
 - Complement the positive number and add 1
- Most architectures use two's complement
- Given n-bit signed integer:
 - Positive range: $0 - 2^{n-1} - 1$
 - Negative range: -2^{n-1}

| type | bytes (32-bit) | bytes (64-bit) | 32-bit range | printf |
|--------------------|----------------|----------------|---------------------------|--------|
| char | 1 | 1 | [0, 255] | %c |
| short int | 2 | 2 | [-32768, 32767] | %hd |
| unsigned short int | 2 | 2 | [0, 65535] | %hu |
| int | 4 | 4 | [-2147483648, 2147483647] | %d |
| unsigned int | 4 | 4 | [0, 4294967295] | %u |
| long int | 4 | 8 | [-2147483648, | %ld |

Beware of integer overflows

- Spot the bug in the following:

```
for (uint32_t n = 10; n >=0; --n) {  
    printf("do I ever terminate?\n");  
}
```

- What does the following print?

```
int x = 0xffffffff;  
printf("%d\n", x);
```



The latest news from Google AI

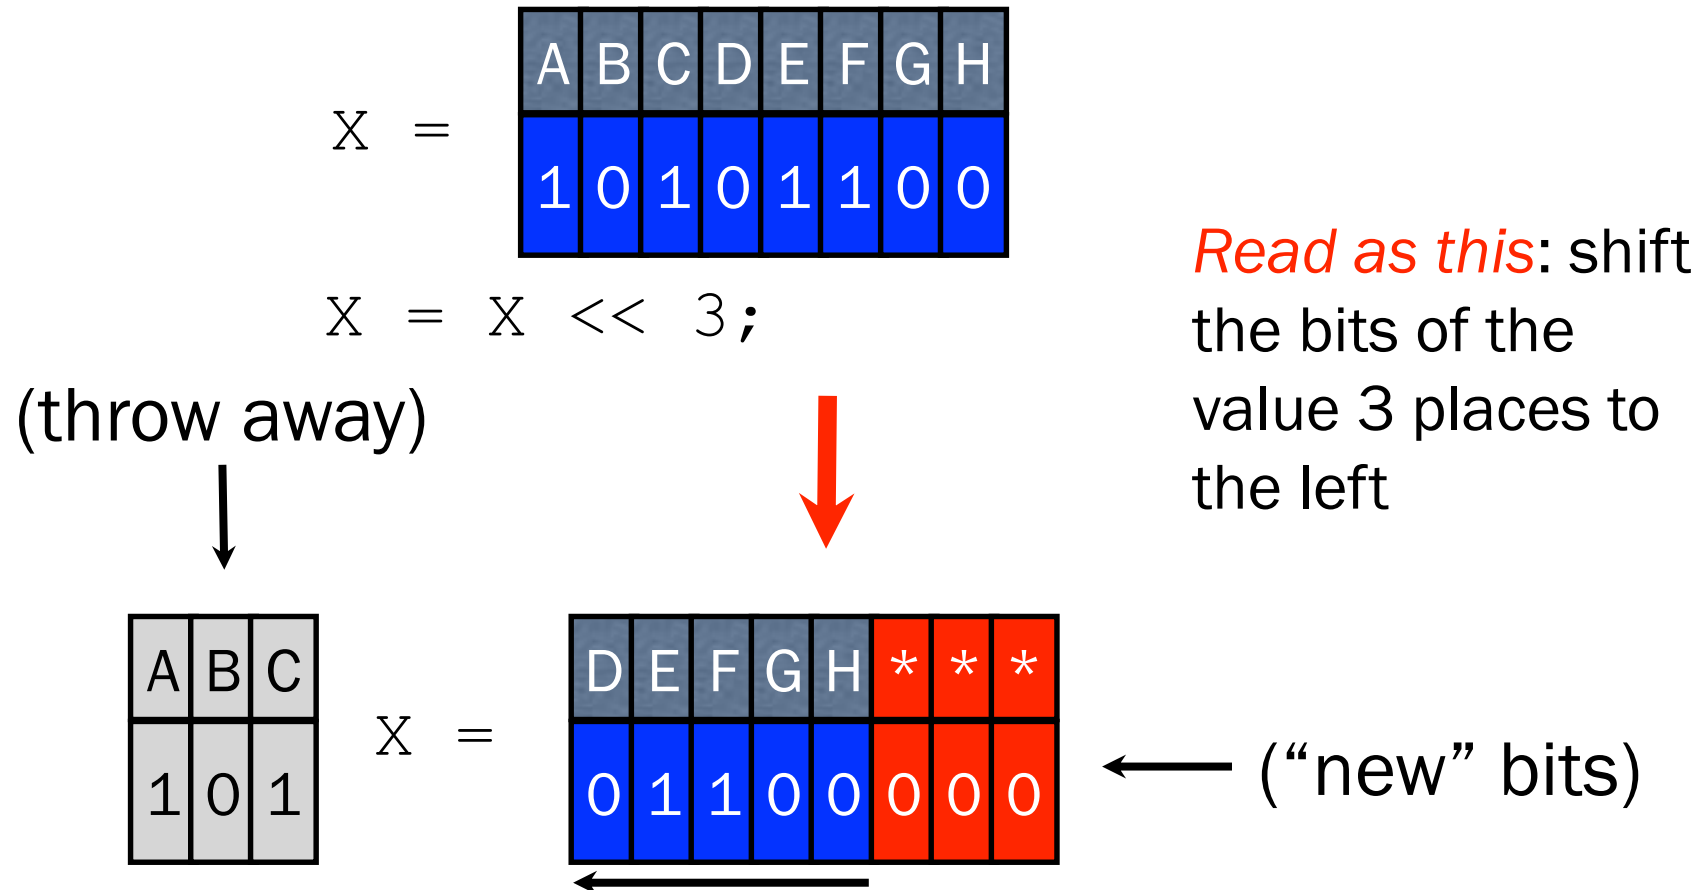
Extra, Extra - Read All About It: Nearly All Binary Searches and Mergesorts are Broken

Friday, June 2, 2006

Posted by Joshua Bloch, Software Engineer

Shift Operations

- A shift operator (\ll or \gg) moves bits to the right or left, throwing away bits and adding bits as necessary



Putting it all together

- Suppose you want to place multiple values in the same 32-bit integer
 - Value **a** in least significant byte
 - Value **b** in 2nd byte
 - Value **c** in 3rd byte
 - Value **d** in 4th byte

| | | | | |
|--------|-------|-------|------|-----|
| Bits | 31-24 | 23-16 | 8-15 | 0-7 |
| Values | d | c | b | a |

Using bit operations ...



PennState

```
uint32_t pack_bytes(uint32_t a, uint32_t b, uint32_t c, uint32_t d) {  
  
    // Setup some local values  
    uint32_t retval = 0x0, tempa, tempb, tempc, tempd;  
  
    tempa = a&0xff; // Make sure you are only getting the bottom 8 bits  
    tempb = (b&0xff) << 8; // Shift value to the second byte  
    tempc = (c&0xff) << 16; // Shift value to the third byte  
    tempd = (d&0xff) << 24; // Shift value to the top byte  
    retval = tempa|tempb|tempc|tempd; // Now combine all of the values  
  
    // Print out all of the values  
    printf("A: 0x%08x\n", tempa);  
    printf("B: 0x%08x\n", tempb);  
    printf("C: 0x%08x\n", tempc);  
    printf("D: 0x%08x\n", tempd);  
  
    // Return the computed value  
    return retval;  
}  
  
...  
  
printf("Packed bytes : 0x%08x\n", pack_bytes(0x111, 0x222, 0x333, 0x444));
```

A: 0x00000011
B: 0x00002200
C: 0x00330000
D: 0x44000000
Packed bytes : 0x44332211