



PennState

CMPS 311 - Introduction to Systems Programming

File Input/Output


Professors:

Suman Saha

(Slides are mostly by

Professor Patrick McDaniel

and *Professor Abutalib Aghayev*)




Input/Output



PennState


- Input/output is the process of moving bytes into and out of devices, files, networks, etc.
 - terminal/keyboard (terminal I/O)
 - secondary storage (file I/O)
 - network (network I/O)
- Different I/O types require different interfaces
 - terminal I/O is messy
 - full of legacy details (not covered)
 - we will cover file I/O and network I/O



Buffered vs. Unbuffered




- When the system is buffering
 - It may read more than requested in the expectation you will read more later (**read buffering**)
 - it may not commit all bytes to the target (**write buffering**)



Blocking vs. Nonblocking vs. Asynchronous PennState

- Non-blocking I/O
 - The call does not wait for the read or write to complete before returning
 - Thus a write/read may commit/return some, all, or none of the data requested
 - When fewer than request bytes are read/written this is called a **short read** or **short write**
- Asynchronous I/O
 - Uses a different API than blocking/non-blocking I/O
 - I/O request returns immediately
 - A callback is generated when I/O completes
- How you program I/O operations is dependent on the blocking behavior of I/O you are using.




Terminal IO



- There are three default terminal channels.

- `STDIN`
- `STDOUT`
- `STDERR`



- UNIX commands/programs for terminal output
 - `echo` - prints out formatted output to terminal `STDOUT`
 - e.g., `echo "hello world"`
 - `cat` - prints out file (or `STDIN`) contents to `STDOUT`
 - e.g., `cat mdadm.c`
 - `less` - provides a read-only viewer for input (or file)
 - e.g., `less mdadm.c`

IO Redirection



PennState

- Redirection uses file for inputs, outputs, or both
 - **Output redirection** sends the output of a program to a file (re-directs to a file), e.g.,
 - `echo "cmpsc311 output redirection" > this.dat`

```
$ echo "cmpsc311 output redirection" > this.dat
$ cat this.dat
cmpsc311 output redirection
```

- **Input redirection** uses the contents of a file as the program input (re-directs from a file), e.g.,
 - `cat < this.dat`

```
$ cat < this.dat
cmpsc311 output redirection
```

- You can also do both at the same time, e.g.,
 - `cat < this.dat > other.dat`

Reading from STDIN vs from a file



PennState

```
#include <stdio.h>

int main(void) {
    char buf[80];
    printf("What is your name? ");
    scanf("%s", buf);
    printf("Hello, %s\n", buf);

    return 0;
}
```

```
$ ./hello
What is your name? Neo
Hello, Neo
$ cat name
Trinity
$ ./hello <name
What is your name? Hello, Trinity
$ ./hello <name >out
$ cat out
What is your name? Hello, Trinity
$
```

Pipes



PennState

- Pipes take the output from one program and uses it as input for another, e.g.,
 - `cat this.dat | less`
- You can also chain pipes together, e.g.,
 - `cat numbers.txt | sort -n | cat`


```
3$ cat numbers.txt
14
21
7
4
$ cat numbers.txt | sort -n | cat
4
7
14
21
$
```


libc



- **libc** is the standard library for the C programming language. It contains the code and interfaces we use for basic program operation and to interact with the parent operating system. Basic interfaces:

- `stdio.h` – declarations for input/output
- `stdlib.h` – declarations for misc system interfaces
- `stdint.h` – declarations for basic integer data types
- `signal.h` – declarations for OS signals and functions
- `math.h` – declarations of many useful math functions
- `time.h` – declarations for basic time handling functions
- ... many, many more



Library call vs system call



- Difference between `open` and `fopen`
 - `open` is a system call

DESCRIPTION

The `open()` system call opens the file specified by `pathname`. If the specified file does not exist, it may optionally (if `O_CREAT` is specified in `flags`) be created by `open()`.

- `fopen` is a library call

DESCRIPTION

The `fopen()` function opens the file whose name is the string pointed to by `pathname` and associates a stream with it.

man man



PennState

- Systems calls, library calls, etc., have their own sections


```
1 Executable programs or shell commands
2 System calls (functions provided by the kernel)
3 Library calls (functions within program libraries)
4 Special files (usually found in /dev)
5 File formats and conventions, e.g. /etc/passwd
6 Games
7 Miscellaneous (including macro packages and conventions), e.g.
  man(7), groff(7)
8 System administration commands (usually only for root)
9 Kernel routines [Non standard]
```

- Use “man 3 fopen” to read the manual page of fopen library call
- You will see references such as “foo(2)” or “bar(3)”
 - the number in parenthesis refers to the manual page section, implies the call type

File IO



- File IO provides random access to a file within the filesystem:
 - With a specific “**path**” (location of the file)
 - At any point in time it has location pointer in the file
 - Next reads and writes will begin at that position
 - All file I/O works in the following way
 - open the file
 - read/write the contents
 - close the file



Locating files for IO



PennState

- An **absolute path** fully specifies the directories and filename itself from the filesystem root “/”, e.g.,

`/home/mcdaniel/courses/cmpsc311-sum19/this.dat`

- A **relative path** is the directories and filename from (or relative to) the current directory, e.g.,

`./courses/cmpsc311-sum19/this.dat`

`courses/cmpsc311-sum19/this.dat`

- All of these references refer to the same file!

FILE* based IO



PennState

- One of the basic ways to manage input and output is to use the FILE set of functions provided by **libc**.
 - The FILE structure is a data structure created to manage input and output for a file
 - An abstraction of “high level” file I/O that avoids some of the details of programming
 - Almost always used for reading and writing ASCII data

```
(gdb) p *file
$3 = {_flags = -72539008, _IO_read_ptr = 0x0, _IO_read_end = 0x0,
      _IO_read_base = 0x0, _IO_write_base = 0x0, _IO_write_ptr = 0x0,
      _IO_write_end = 0x0, _IO_buf_base = 0x0, _IO_buf_end = 0x0,
      _IO_save_base = 0x0, _IO_backup_base = 0x0, _IO_save_end = 0x0,
      _markers = 0x0, _chain = 0x7ffff7dd41a0 <_IO_2_1_stderr_>, _fileno =
      7, _flags2 = 0, _old_offset = 0, _cur_column = 0,
      _vtable_offset = 0 '\000', _shortbuf = "", _lock = 0x6020f0, _offset
      = -1, __pad1 = 0x0, __pad2 = 0x602100, __pad3 = 0x0, __pad4 = 0x0,
      __pad5 = 0, _mode = 0, _unused2 = '\000' <repeats 19 times>}
```

fopen()



PennState

- The fopen function opens a file for IO and returns a pointer to a FILE* structure:

```
FILE *fopen(const char *path, const char *mode);
```

- Where,
 - `path` is a string containing the absolute or relative path to the file to be opened.
 - `mode` is a string describing the ways the file will be used
 - For example, `FILE *file = fopen(filename, "r+");`
 - Returns a pointer to FILE* if successful, NULL otherwise
 - You don't have to allocate or deallocate the FILE* structure

fopen()



PennState

- The fopen function opens a file for IO and returns a pointer to a FILE* structure:

```
FILE *fopen(const char *path, const char *mode);
```

- Where,


- `path` is a string containing the file to be opened.
 - `mode` is a string describing the file to be opened.
- A FILE* structure is also referred to as a *stream*.
- For example, `FILE *file = fopen(filename, "r+");`
 - Returns a pointer to FILE* if successful, NULL otherwise
 - You don't have to allocate or deallocate the FILE* structure

fopen modes



PennState

- “**r**” - Open text file for reading. The stream is positioned at the beginning of the file.
- “**r+**” - Open for reading and writing. The stream is positioned at the beginning of the file.
- “**w**” - Truncate file to zero length or create text file for writing. The stream is positioned at the beginning of the file.
- “**w+**” - Open for reading and writing. The file is created if it does not exist, otherwise it is truncated.
- “**a**” - Open for appending (writing at end of file). The file is created if it does not exist.
- “**a+**” - Open for reading and appending (writing at end of file). The file is created if it does not exist.



Reading the file



PennState

- There are two dominant ways to read the file, `fscanf` and `fgets`
 - `fscanf` reads the data from the file just like `scanf`, just reading and writing, e.g.,

```
if ( fscanf( file, "%d %d %d\n", &x, &y, &z ) == 3 ) {  
    printf( "Read coordinates [%d,%d,%d]\n", x, y, z );  
}
```

- `fgets` reads a line of text from the file, e.g.,

```
if ( fgets(str,128,file) != NULL ) {  
    printf( "Read line [%s]\n", str );  
}
```

Writing the file



- There are two dominant ways to write the file, `fprintf` and `fputs`
 - `fprintf` writes the data to the file just like `printf`, just reading and writing, e.g.,

```
fprintf( file, "%d %d %d\n", x, y, z );
```

- `fputs` writes a line of text to the file, e.g.,

```
if ( fputs(str,file) != NULL ) {  
    printf( "wrote line [%s]\n", str );  
}
```

fflush



PennState

- `FILE*`-based IO is buffered
- `fflush` attempts to reset/the flush state

```
int fflush(FILE *stream);
```

- `FILE*`-based writes are buffered, so there may be data written, but not yet pushed to the OS.
 - `fflush()` forces a write of all buffered data
- `FILE*`-based reads are buffered, so the current data (in the process space) may not be current
 - `fflush()` discards buffered data from the underlying file
- If the stream argument is `NULL`, `fflush()` flushes **all** open output streams
- `fflush()` does not guarantee that data is safely on disk; for that use `fsync(2)`

fclose()



- `fclose()` closes the file and releases the memory associated with the `FILE*` structure.

```
fclose( file );  
file = NULL;
```

Note: `fclose()` implicitly flushes the data to OS.

Putting it all together ...



PennState

```
int show_fopen( void ) {

    int x, y, z;
    FILE *file;
    char *filename = "/tmp/fopen.dat", str[128];
    file = fopen( filename, "r+" );

    // open for reading and writing
    if ( file == NULL ) {
        fprintf( stderr, "fopen() failed, error=%s\n", strerror(errno) );
        return( -1 );
    }

    // Read until you reach the end
    while ( !feof(file) ) {
        if ( fscanf( file, "%d %d %d\n", &x, &y, &z ) == 3 ) {
            printf( "Read coordinates [%d,%d,%d]\n", x, y, z );
        }
        if ( !feof(file) ) {
            fgets(str,128,file); // Need to get end of previous line
            if ( fgets(str,128,file) != NULL ) {
                printf( "Read line [%s]\n", str );
            }
        }
    }
}
```

Putting it all together ...



PennState

```
// Now add some new coordinates
x = 21;
y = 34;
z = 98;
fprintf( file, "%d %d %d\n", x, y, z );
printf( "Wrote %d %d %d\n", x, y, z );
if ( fputs(str,file) >= 0 ) {
    printf( "wrote line [%s]\n", str );
}
fflush( file );

// Close the file and return
fclose( file );
return( 0 );
}
```

```
$ cat /tmp/fopen.dat
1 2 3
4 5 6
11 12 14
16 17 23
$ ./io
This is cmpsc311, IO example
Read coordinates [1,2,3]
Read line [11 12 14
]
Read coordinates [16,17,23]
Wrote 21 34 98
wrote line [11 12 14
]
$ cat /tmp/fopen.dat
1 2 3
4 5 6
11 12 14
16 17 23
21 34 98
11 12 14
$
```

open ()



- The open system call opens a file for IO and returns an integer file handle:

```
int open(const char *path, int flags, mode_t mode);
```

- Where,
 - `path` is a string containing the absolute or relative path to the file to be opened.
 - `flags` indicates the kind of open you are requesting
 - `mode` sets a security policy for the file

- `open ()` returns a “file handle”

open() flags



- The “flags” to open with
 - `O_RDONLY` - read only
 - `O_WRONLY` - write only
 - `O_RDWR` - read and write
- Options
 - `O_CREAT` - If the file does not exist it will be created.
 - `O_EXCL` Ensure that this call creates the file, fail otherwise (fail if already exists)
 - `O_TRUNC` - If the file already exists it will be truncated to length 0.



Note: You bitwise or (`|`) the mode/options you want

Access Control in UNIX




- The UNIX filesystem implements **discretionary access control** through file permissions set by user
 - The permissions are set at the discretion of the user
- Every file in the file system has a set of bits which determine who has access to the files
 - **User** - the owner is typically the creator of the file, and the entity in control of the access control policy
 - **Group** - a set of users on the system setup by the admin
 - **World** - the set of everyone on the system
- Note: this can be overridden by the “**root**” user

UNIX filesystem rights ...



- There are three rights in the UNIX filesystem
 - **READ** - allows the subject (process) to read the contents of the file.
 - **WRITE** - allows the subject (process) to alter the contents of the file.
 - **EXECUTE** - allows the subject (process) to execute the contents of the file (e.g., shell program, executable, ...)




- **Q:** why is execute a right?
- **Q:** does read implicitly give you the right to execute?

UNIX Access Policy



- Really, this is a bit string encoding an access policy:



- And a policy is encoded as “**r**”, “**w**”, “**x**” if enabled, and “**-**” if not, e.g,

rwXrw---x

- Says user can read, write and execute, group can read and write, and world can execute only.

UNIX Access Policy



PennState

- Really, this is a bit string encoding an access policy:

`rwX rwX rwX`

```
$ ls -l .
total 20
-rw-rw-r-- 1 mcdaniel professor 12 Oct 10 14:18 fopen.dat
-rwxrwxr-x 1 mcdaniel professor 12058 Oct 10 15:42 io
-rw-rw-r-- 1 mcdaniel professor 1176 Oct 10 15:42 io.c
$
```

- And a policy, is encoded as `rwXrw---x`, and “-” if not, e.g,

`rwXrw---x`

- Says user can read, write and execute, group can read and write, and world can execute only.

Setting an access policy



PennState

- Specify a file access policy by bit-wise ORing (|):
 - **S_IRWXU** **00700** user (file owner) has read, write and execute
 - **S_IRUSR** **00400** user has read permission
 - **S_IWUSR** **00200** user has write permission
 - **S_IXUSR** **00100** user has execute permission
 - **S_IRWXG** **00070** group has read, write and execute permission
 - **S_IRGRP** **00040** group has read permission
 - **S_IWGRP** **00020** group has write permission
 - **S_IXGRP** **00010** group has execute permission
 - **S_IRWXO** **00007** world has read, write and execute permission
 - **S_IROTH** **00004** world has read permission
 - **S_IWOTH** **00002** world has write permission
 - **S_IXOTH** **00001** world has execute permission

Putting it together ...



PennState

- So an open looks something like ...

```
// Setup the file for creating and open
flags = O_WRONLY|O_CREAT|O_EXCL; // Create a NEW file (no overwrite)
mode = S_IRUSR|S_IWUSR|S_IRGRP;   // User can read/write, group read
fhandle = open( filename, flags, mode );
if ( fhandle == -1 ) {
    fprintf( stderr, "open() failed, error=%s\n", strerror(errno) );
    return( -1 );
}
```


Q: But why is an `int` returned by `open()` a file?

File descriptor



PennState

- A **file descriptor** is an index assigned by the kernel into a table of file information maintained in the OS
 - The file descriptor table is unique to each process and contains the details of open files.
 - File descriptors are used to reference when calling the I/O system calls.
 - The kernel accesses the file for the process and returns the results in system call response.



Reading and Writing



- Primitive reading and writing mechanisms that only process only blocks of opaque data:

```
ssize_t write(int fd, const void *buf, size_t count);  
ssize_t read(int fd, void *buf, size_t count);
```

- Where `fd` is the file descriptor, `buf` is an array of bytes to write from or read into, and `count` is the number of bytes to read or write
- Both `read()` and `write()` returned the number of bytes read and written.
 - Be sure to always check the result (!!!)
- On reads, you are responsible for supplying a buffer that is large enough to put the output into.

close()



- `close()` closes the file and deletes the file's entry in the file descriptor table

```
close(fhandle);  
fhandle = -1;
```

Note: Always reset your file handles to `-1` to avoid use after close.

Putting it all together ...



PennState

```
int show_open( void ) {

    // Setup variables
    char *filename = "/tmp/open.dat";
    int vals[1000] = { [0 ... 999] = 0xff }, vals2[1000];
    int fhandle, flags;
    mode_t mode;

    // Setup the file for creating and open
    flags = O_WRONLY|O_CREAT|O_EXCL; // Create a NEW file (no overwrite)
    mode = S_IRUSR|S_IWUSR|S_IRGRP;   // User can read/write, group read
    fhandle = open( filename, flags, mode );
    if ( fhandle == -1 ) {
        fprintf( stderr, "open() failed, error=%s\n", strerror(errno) );
        return( -1 );
    }

    // Now write the array to the file
    if ( write(fhandle, (char *)vals, sizeof(vals)) != sizeof(vals) ) {
        fprintf( stderr, "write() failed, error=%s\n", strerror(errno) );
        return( -1 );
    }
    close( fhandle );
    fhandle = -1;
}
```

Putting it all together ...



PennState

```
// Setup the file for reading
flags = O_RDONLY; // Read an existing file
fhandle = open( filename, flags, 0 );
if ( fhandle == -1 ) {
    fprintf( stderr, "open() failed, error=%s\n", strerror(errno) );
    return( -1 );
}

// Now read the array from the file
if ( read(fhandle, (char *)vals2, sizeof(vals2)) != sizeof(vals2) ) {
    fprintf( stderr, "read() failed, error=%s\n", strerror(errno) );
    return( -1 );
}
close( fhandle );
return( 0 );
}
```

```
$ ./io
$ $ od -x -N 256 /tmp/open.dat
0000000 00ff 0000 00ff 0000 00ff 0000 00ff 0000
*
0000400
```

fopen() vs. open()



- Key differences between `fopen` and `open`
 - `fopen` provides you with buffering IO that may or may not turn out to be a faster than what you're doing with `open`.
 - `fopen` does line ending translation if the file is not opened in **binary mode**, which can be very helpful if your program is ever ported to a non-Unix environment.
 - A `FILE *` gives you the ability to use `fscanf` and other `stdio` functions that parse out data and support formatted output.
- IMO: use `FILE*` style I/O for ASCII processing, and file handle I/O for binary data processing.

A parting note ...



PennState

- Each of the styles of I/O requires a different set of include files
 - libc file I/O (i.e. using FILE*, fopen, fclose, ...) requires:

```
#include <stdio.h>
```

- system call file I/O requires:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
```

- Read the manual page to find out what to include for what call

NAME

open, openat, creat - open and possibly create a file

SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```